COMP34112
Natural Language Systems

Coursework 1

In this coursework, you will explore corpus-based approaches to POS tagging and distributional semantics. You may want to use functions that are available in the NLTK framework to implement the features required.

All the datasets needed for the coursework are in Blackboard.

The coursework is worth a total of **10 marks** of the final COMP34412 final mark.

Tasks

1) Part-of-speech (POS) tagging [3 marks]

a) Use a POS-tagger of your choice (e.g. NLTK tagger, Stanford, TreeTagger etc. – see Homework 2) to tag corpus A (see Blackboard).

b) Let's assume that the POS tagging you have generated for the corpus is a gold standard. Use the annotated corpus to estimate word likelihood and tag transition probabilities you would need to be able to disambiguate which of the following two POS tagging results is more likely

(1) *People/NNS continue/VB to/TO inquire/VB the/DT reason/NN for/IN the/DT race/NN for/IN outer/JJ space/NN*

(2) *People/NNS continue/VB to/TO inquire/VB the/DT reason/NN for/IN the/DT race/VB for/IN outer/JJ space/NN*

Explain what you have done and comment/explain the results in the report (½ page).

2) Distributional semantics [7 marks]

a) Implement a program to cluster a given list of target words into n groups based on their distributional co-occurrence patterns. You may first want to construct a word-by-word m X N matrix (m = number of target words; N = number of words in the vocabulary) that captures co-occurrence patterns of the given target words using a given corpus. Define your context and features flexibly so that you can analyse the results in part c below. Your program should take as input a list of words to cluster and a number of clusters. Use any available machine-learning framework for clustering (e.g. sckikit-learn (http://scikit-learn.org/) or Weka (http://www.cs.waikato.ac.nz/ml/weka/)). Explain briefly what you have done in the report (½ page).

b) Use corpus B (available in Blackboard) and target list D (with 50 words, also available in Blackboard) to evaluate the results of your clustering. Use the following *pseudoword* disambiguation approach: for each target word, randomly substitute half of its occurrences in the corpus with its reverse (e.g., "*procedure*" will be transformed into "*erudecorp*"). Now, apply your clustering algorithm to the list of 100 target words, which contains original words and their reverses, producing 50 clusters. If you generate 50 clusters, how many of them will contain "correct" pairs (i.e., a word and its reverse)? Repeat this process 5 times and give the average accuracy. Explain the result in the report (½ page).

c) Analyse the impact of (1) the *size of context*, (2) *type of features* and (3) *training data* on the quality of generated clusters. To analyse the contribution of contextual representation, consider different ways of constructing a word-by-word matrix (i.e. vary the dimensions of the context window) and experiment with different

definitions of context (stems vs. words). To analyse the impact of training data, in addition to corpus B, use also corpus C (available in Blackboard) and "train" your system on each corpora separately, and also on their combination. Comment the results and report any difference (1 page). What other type(s) of feature you may consider using (you don't need to implement or analyse the impact of the additionally proposed feature(s))?

Submissions

The deadline for submissions is **6pm on Friday February 28th 2020**. Your submission should be a zip file uploaded via Blackboard. For each task you should submit a write-up (in a single pdf) that clearly explains what you have done and presents the results. You should also submit your source code and the output of your code (where applicable). The README file should clearly specify how to run your program.

Data

All the datasets are available in Blackboard.